

SSE3044 Introduction to Operating Systems

Prof. Jinkyu Jeong

Project 2. CPU scheduling

TA)

송원석(wonsuk.song@csi.skku.edu)

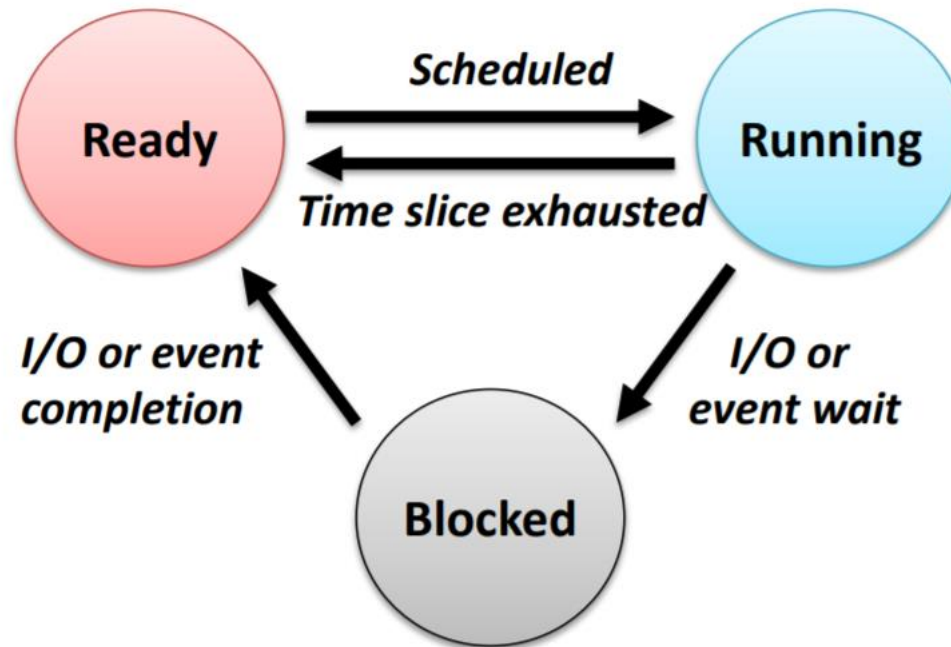
진승우(seungwoo.jin@csi.skku.edu)

Project plan

- Total 6 projects
 - 0) Booting xv6 operating system
 - 1) System call
 - 2) CPU scheduling
 - Linux CFS scheduler
 - 3) Virtual memory
 - 4) Page replacement
 - 5) File systems

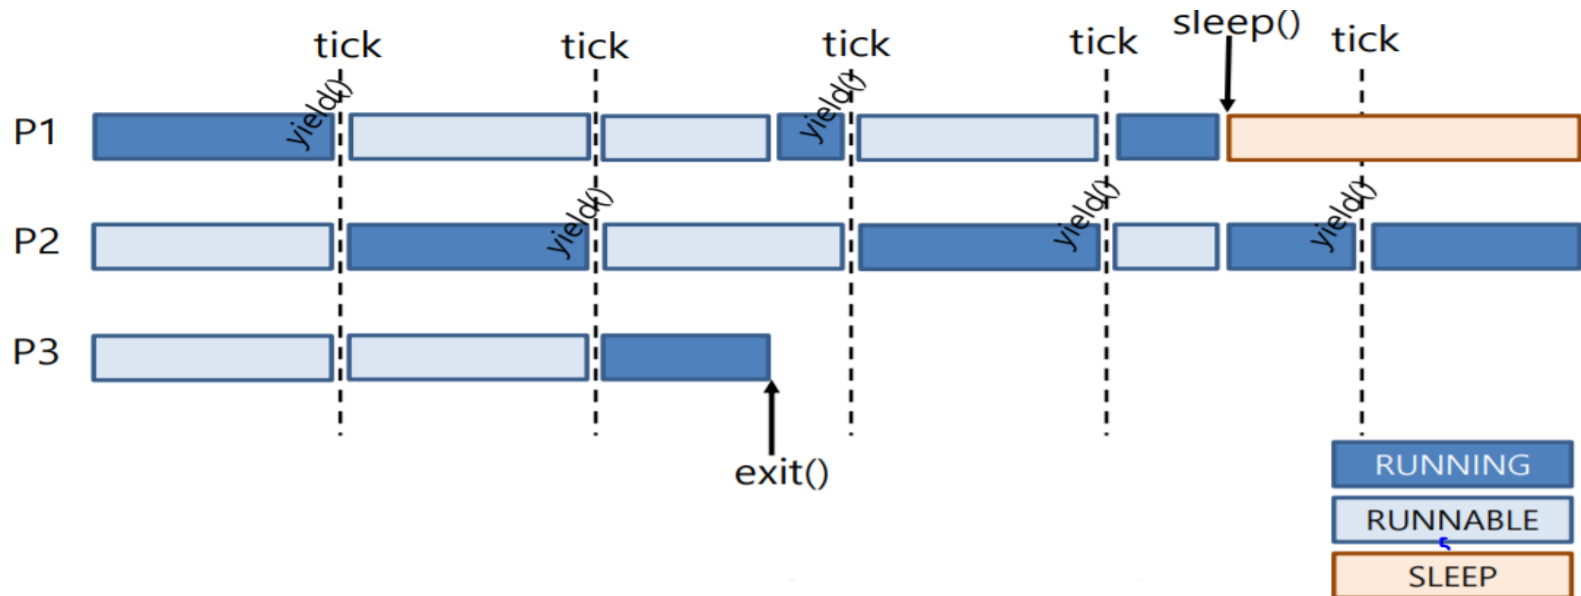
CPU scheduling

- Selects from among the processes in memory that are ready to execute, and allocates CPU to one of them



How current scheduler works in xv6?

- Every timer IRQ enforces a yield of a CPU
- Process to be scheduled to RUNNING will be chosen in round-robin manner



Strawman scheduler

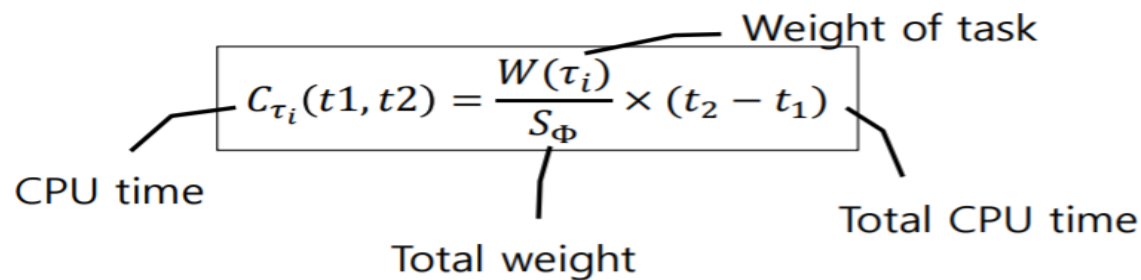
- Organize all processes as a simple list
- In `schedule()`:
 - Pick first one on list to run next
 - Put suspended task at the end of the list
- Problems?
 - Allows only round-robin scheduling
 - Can't prioritize tasks

Fair scheduling

- And, how should time slices be distributed according to priority?
 - The difference of time slice by the nice value is not fair
 - E.g, processes with nice value 0 and 1 are given 100ms and 95ms
 - Processes with nice value 18 and 19 are given 10ms and 5ms
 - The differences are same to 5ms, but it's not proportional
 - To solve this problem, From the Linux kernel 2.6.23, CFS(Completely Fair Scheduler) was used

CFS (Completely Fair Scheduling)

- Linux default scheduler
- Basic concept
 - The CPU is allocated to the process in proportion to its weight
 - CPU time of any task satisfies in any given time between t_1 and t_2



The diagram shows the formula for CFS CPU time allocation: $C_{\tau_i}(t_1, t_2) = \frac{W(\tau_i)}{S_{\Phi}} \times (t_2 - t_1)$. Hand-drawn lines point from text labels to parts of the formula: 'CPU time' points to $C_{\tau_i}(t_1, t_2)$, 'Weight of task' points to $W(\tau_i)$, 'Total weight' points to S_{Φ} , and 'Total CPU time' points to $(t_2 - t_1)$.

$$C_{\tau_i}(t_1, t_2) = \frac{W(\tau_i)}{S_{\Phi}} \times (t_2 - t_1)$$

- Nice to weight
 - Difference in nice by 1 provides 10% more (or less) CPU time
 - However, the larger the absolute value of nice, the smaller the ratio between the two values
 - Therefore, a new concept “weight”
 - Although there is formula, use pre-defined as array in Linux

$$Weight = 1024(weight\ of\ nice\ 0) \times (1.24)^{-NICE}$$

CFS parameters

- Time slice

- Task's minimum time to be executed before it is preempted
- Allocated to the process in proportion to its weight

$$\text{time slice} = \text{scheduling_latency} * \frac{\text{weight of task}}{\text{total weight of runqueue}}$$

- Scheduling latency (6ms by default)
 - Minimum time period to satisfy proportional CPU time distribution

- vruntime (virtual runtime)

- accounts for how long a process has run proportional to its weight
- It's easy to compare how fairly the CPU is allocated
- By comparing this value, allowing you to select the next process to be performed

$$\text{vruntime} = \text{actual runtime} \times \frac{\text{nice 0's weight}}{\text{weight of task}}$$

CFS scheduling

1. A task with minimum virtual runtime is scheduled
2. Scheduled task gets time slice proportional to its weight/total weight
3. During task running, virtual runtime is updated
4. After task runs more than time slice, go back to 1

Project 2. Implement CFS on xv6

- Implement CFS on xv6

- Select process with minimum virtual runtime among runnable processes
- Update runtime, vruntime every timer interrupt
- If task runs more than time slice, enforces an yield of the CPU
- Default nice value is 0, range from -5 to 5, and weight of nice 0 is 1024

- Nice(-5~5) to weight(Although there is formula, use pre-defined as array like Linux)

$$Weight = 1024(weight\ of\ nice\ 0) \times (1.24)^{-NICE}$$

/* -5 */	3121,	2501,	1991,	1586,	1277,
/* 0 */	1024,	820,	655,	526,	423,
/* 5 */	335,				

- Time slice calculation (our scheduling latency is 10ticks)

$$time\ slice = 10tick \times \frac{weight\ of\ current\ process}{total\ weight\ of\ runnable\ processes}$$

- vruntime calculation

$$vruntime+ = \Delta\ runtime * \frac{weight\ of\ nice0(1024)}{weight\ of\ current\ process}$$

Project 2. Implement CFS on xv6

- How about forked process?
 - A process inherits the parent process's vruntime
- How about woken process?
 - When a process is woken up, it's virtual runtime gets (minimum vruntime of processes in the ready queue – vruntime(1tick))
$$vruntime(1tick) = 1tick * \frac{\text{weight of nice0}(1024)}{\text{weight of current process}}$$
- Do NOT call sched() during a wake-up of a process
 - Ensuring that the time slice of the current process expires
 - Woken-up process will have the minimum vruntime (due to the formula above)
 - But, we do NOT want to schedule the woken-up process before the time slice of current process expires.
 - This is by default in xv6

Project 2. Implement CFS on xv6

- To check that CFS is implemented properly, ps() should be modified
- Expected output (testcfs.c)

```
$ testcfs
=== TEST START ===
name      pid    state      priority    runtime/weight  runtime    vruntime    tick 3513000
init      1      SLEEPING    5           2             1000      1000
sh        2      SLEEPING    5           0             0         0
testcfs   3      RUNNABLE    5           1764          591000    589000
testcfs   4      RUNNING     0           1757          1800000    588600
```

- Print out the following information about the processes
- Use millitick unit (multiply the tick by 1000)
 - runtime, vruntime, total tick
 - Do NOT use float/double types to present runtime and vruntime
 - Kernel avoid floating point operation as much as possible
- Indents of name section should be aligned even if process has long name (up to 10 letters) or very large value... (runtime, vruntime)

Project 2. Implementation details

- Project 2 should be done based on your project 1 code
- Consider the case of integer overflow vruntime
 - Even if over the scope of integer, shall operate without problems
 - And it has to be printed normally
 - Do not worry about runtime, total tick
- FAQ
 - Q : My time slice is 6.5. However, what if timer interrupt occurs every 1 tick?
(context switch can occur only with 1 tick)
 - A : Tasks will run over its time slice (7 ticks) & add vruntime
- Project 2 is very complex to implement, so please ask a lot of questions!

Submission

- Please implement CFS on xv6
- Compress your source code and upload on i-Campus
- How to compress your project
 - If you insert the user program, Modify the 'EXTRA' in Makefile
 - make dist
 - make tar
 - then, tar.gz file will be generated automatically
 - Rename to studentID-project2.tar.gz and upload in email
- Submit a report together, the file format of the report is limited to pdf
 - There is no limit to the format of the contents
 - But, include your description of your code and execution screen

Submission

- File format
 - StudentID-project2.tar.gz
 - StudentID-project2.pdf
- PLEASE DO NOT COPY
 - YOU WILL GET F GRADE IF YOU COPIED
- Due date: 4/21 (Tue.), 23:59:59 PM
 - -25% per day for delayed submission

Questions

- If you have questions, please ask in Piazza
- You can also visit Semiconductor Building #400509
 - Please e-mail TA before visiting
- Reading xv6 commentary will help you a lot
 - <http://csl.skku.edu/uploads/SSE3044S20/book-rev11.pdf>